

1 Race-based disparities in allocation of academic disciplinary actions are associated with
2 county-level rates of bias

3 Travis Riddle¹ & Stacey Sinclair¹

4 ¹ Princeton University

5 Author Note

6 Enter author note here.

7 Correspondence concerning this article should be addressed to Travis Riddle. E-mail:

8 triddle@princeton.edu

Abstract

9

10 Enter abstract here.

11 *Keywords:* keywords

12 Word count: X

Race-based disparities in allocation of academic disciplinary actions are associated with county-level rates of bias

Introduction

Methods

Data Sources

We used three distinct data sources for the work described here. The academic disciplinary data is part of the Civil Rights Data Collection (CRDC) through the US Department of Education. The dataset we used comes from the 2013-2014 academic year and has data on “all public local and educational agencies and schools, including long-term secure juvenile justice facilities, charter schools, alternative schools, and schools serving students with disabilities.” In total, the CRDC data represents 95507 institutions enrolling approximately 50 million students, of which approximately 25 million are white and 7.8 million are Black¹. Previous work using these data have identified a number of districts whose data are in error, and have excluded juvenile justice facilities, as these institutions constitute dramatically different educational environments, where the meaning of disciplinary actions may be quite different (Losen et al., 2015). After these exclusions are applied, the final sample used for modeling consists of 90002 institutions, enrolling 32 million black or white students, of which 25 million are white and 7 million are black.

We obtained county-level demographic information for use as covariates and state-level demographic information for use in post-stratification from the US Census Bureau. For both county-level and state-level demographics, we use 5-year estimates for the period ending in 2014 from the American Community Survey, which surveys around 295k households per

¹We note that there are a number of differences between the analyses we registered and those presented in the main text. Our general conclusions are largely the same for both sets of analyses. We opted to report the modified analyses for reasons of clarity and to remain congruent with previous research on the same topics. The registered analyses can be found, in full, in the appendix.

month.

We used measurements of implicit and explicit bias available from data collected through Project Implicit (Xu, Nosek, & Greenwald, 2014). From this data, we used only respondents who had geographic information that would allow us to place them in a United States county, identified as White, and visited the site before 2015. This consisted of approximately 1.1 million total respondents from 3091 counties.

A subset of years in the Project Implicit data also collected occupational information from respondents. As identified in our pre-analysis plan, we took advantage of the presence of primary and secondary educators in these data to test whether any associations between bias and race-based differences in the rates of disciplinary action were stronger among these respondents. Filtering for only white individuals who identified as primary, secondary, special education, and other teachers and instructors (occupation codes 25-2000 and 25-3000) reduced the dataset to 63552 respondents. In order to assure that our estimates were reasonably stable, we limited analysis to only counties that had at least 50 respondents. As such, our teacher analysis is limited to just 287 counties. Additionally, because we do not know of any state-level demographic estimates for teachers, we are unable to perform post-stratification for these data.

Measures

The primary outcome in this analysis is a count of the number of students by race (black and white) who received one of several types of disciplinary action. We report here rates of out-of-school suspension, in-school suspension, school-related arrests, law enforcement referrals, and total number of expulsions of any type.

For both the post-stratification procedure (described below) and the actual statistical models used for inference, we used the same set of covariates. These covariates were at the state-level for post-stratification and were at the county-level for the statistical models used for final estimation and inference. Specifically, our population based covariates were the total

population count, the proportion of the population that is black, the proportion of the population that is white and the ratio of black-to-white persons. We also used socioeconomic covariates. We used the percent of individuals aged 16 or over who were in the labor force but unemployed, the median household income, and the percentage of all families whose income is below the poverty line.

Finally, for the implicit and explicit bias measures, we relied on two primary variables from the Project Implicit data. Implicit bias was assessed via an Implicit Association Test (IAT). This test uses a speeded dual-categorization task in which individuals must quickly categorize black and white faces and “good” and “bad” words with key presses. The difference in how quickly and accurately participants are able to pair white faces with “good” words and black faces with “bad” words in comparison to the inverse is thought to reflect implicit associations between the two races and positive and negative affective reactions. This association is indicated in the IAT D-score, which we used as a measure of implicit bias. Our measure of explicit bias is the difference between reported warmth towards whites (i.e. *how warm or cold do you feel towards Whites? 0=very cold, 10=very warm*) and reported warmth toward blacks.

Data analysis

Data analysis proceeded in two steps. We first estimated county-level implicit and explicit bias using multilevel regression and post-stratification. Post-stratification is a valuable procedure in obtaining accurate geographical population-based estimates because it allows a non-representative sample (e.g Project Implicit) to more closely resemble the true population, and it regularizes extreme observations with little data to support them (e.g. a county with only a handful of respondents with especially high or low scores) (Gelman & Little, 1997; Park, Gelman, & Bafumi, 2004). Following past work (Leitner, Hehman, Ayduk, & Mendoza-Denton, 2016), we identified age as one dimension along which IAT respondents differed from the general population in ways that could bias our conclusions (Gonsalkorale,

Sherman, & Klauer, 2009). Our post-stratification weighting scheme is as follows: We first grouped respondents into five age group categories (15-24, 25-34, 35-54, 55-75, and 75+). We next fit multilevel models estimating bias (implicit and explicit biases separately) as a function of our state-level covariates (the “fixed” effects), and allowed the estimates to vary by age bin, county, and state (the “random” effects). Next, we determined the population of whites in each county in these age groups using the American Community Survey’s 5-year estimates ending in 2014. Finally, we used our estimated models to predict the expected response for each age bin, in each county. Our final county-level estimates are the average of the values predicted for the 5 age bins, weighted by the population size of that bin in that county. As a result of this procedure, we can be confident that our county-level estimates should more closely approximate what our estimates would look like if the Project Implicit data were truly representative along the age dimension in all counties.

After obtaining these estimates, we use them as predictors in bayesian multilevel logistic regressions. Formally, the likelihood for a given observation is written as a binomial function:

$$\binom{n}{y} \pi^y (1 - \pi)^{n-y}$$

Where y is the observed count of incidents (e.g. number of black or white students suspended), n is the number of at-risk students (e.g. total number of black or white students), and $\pi = g^{-1}(\eta)$ is the probability of the incident occurring. For this analysis, the linear predictor takes the form of a multilevel model with a set of effects that vary over county:

$$\eta = \alpha + X\beta + \gamma_{county}$$

Where α is an intercept that is constant across observations, β represents a set of effects that are also constant across observations (i.e. fixed effects), and γ_{county} represents intercepts and effects of ethnicity that vary across the counties (i.e. random effects).

In addition to the covariates described above, we also include effects of race, implicit

bias, explicit bias, and the two-way interactions between implicit bias and race and explicit bias and race. We fit separate models for each of the outcomes.

Because of the computational demands of fitting such a high-dimensional model to such a large dataset (the full model for each metric would consist of over 6k parameters to approximately 170k observations), we used a consensus monte carlo algorithm to obtain approximate posterior distributions for the parameters of interest (Scott et al., 2016). The approximate posteriors derived from this algorithm have been shown to be nearly indistinguishable from the true posterior, a result we verified using a small subset of our own data.

All numerical predictor variables were standardized at the appropriate level (county, state) before model estimation to help with estimation efficiency and interpretability. We set priors for the intercept and coefficients in the bayesian models to be weakly informative normal distributions centered on zero with a standard deviation of five. All other parameters were left to default values. Data analysis was done in R (R Core Team, 2016) version 3.3.2 running under OS X 10.11.6. Post-stratification was done with lme4, version 1.1.14 (Bates, Mächler, Bolker, & Walker, 2015). Final model fitting was done on the university cluster running Springdale Linux, release 6.9 using rstanarm, version 2.17.2 (Stan Development Team, 2016). We used the implementation of the consensus monte carlo algorithm found in parallelMCMCcombine, version 1.0 (Miroshnikov & Conlon, 2014). Figures were made with ggplot2, version 2.2.1 (Wickham, 2009), with data manipulation done using dplyr version 0.7.2 (Wickham, Francois, Henry, & Müller, 2017) and tidyr, version 0.7.1 (Wickham & Henry, 2017). A full report of session information can be found on the OSF page (...)

Results

Project Implicit Estimates

We first report the results of estimating the implicit and explicit biases in from Project Implicit data. Overall, the individuals in project implicit show a pro-white bias in both

Table 1

Count of students by race receiving each type of disciplinary action

group	metric	students
black	expulsions	36,755.00
black	school arrests	21,456.00
black	in-school suspension	829,706.00
black	law enforcement referral	53,306.00
black	out-of-school suspension	1,850,492.00
white	expulsions	55,832.00
white	school arrests	21,420.00
white	in-school suspension	1,054,172.00
white	law enforcement referral	81,565.00
white	out-of-school suspension	1,768,326.00

implicit ($mean = 0.40$, $sd = 0.41$), and explicit measures ($mean = 0.88$ $sd = 1.83$). When aggregated at the county level and adjusted with poststratification, the summary statistics across counties are similar in terms of their location, but as expected, the variability is much diminished ($mean_{implicit} = 0.40$, $sd_{implicit} = 0.02$; $mean_{explicit} = 0.79$, $sd_{explicit} = 0.15$, where on both scales 0 = no bias, and positive numbers indicate a pro-white bias).

Disciplinary action frequency

Table 1 shows the number of students of each race who were reported having received each of the actions under consideration. The counts range from a low of just 21420 white students arrested to a high of 1850492 black students receiving an out-of-school suspension. Considering the vast differences in the overall number of black and white students, this

simple count already illustrates that black students are disciplined at rates far higher than their white counterparts.

Associations across county

Figure 1 shows the estimate of primary interest for each of the models. The estimates displayed are the coefficients for the interaction between race and each of the two bias measurements. Given that African Americans are the baseline group, negative values for this coefficient indicate that as the bias in a county increases, the gap between the probability of a black student being disciplined and the probability of a white student being disciplined grows.

Several patterns are apparent from this figure. First, with the exceptions of implicit bias and expulsions and law enforcement referrals, all estimates are directionally consistent with higher levels of bias leading to larger differences between groups. Second, these effects are especially consistent for the two types of suspensions. The largest effect estimated is between implicit bias and out-of-school suspensions. The difference in the slope of the association between implicit bias and the log of the odds for out-of-school suspensions between white and black students is estimated to be -0.25, with 95% of the posterior distribution between -0.31 and -0.20 and a proportion $>.99$ of the posterior distribution consistent with a negative effect.

Although not nearly as large of a difference, we have similar certainty with respect to the association between out-of-school suspensions and explicit bias. The difference in the slope of the association between explicit bias and the log of the odds for out-of-school suspensions between white and black students is estimated to be -0.08, with 95% of the posterior distribution between -0.14 and -0.03 and a proportion $>.99$ of the posterior distribution consistent with a negative effect.

The estimated associations for in-school suspensions are smaller still, but are generally consistent with effects of the same direction. For implicit bias, the relevant parameter is

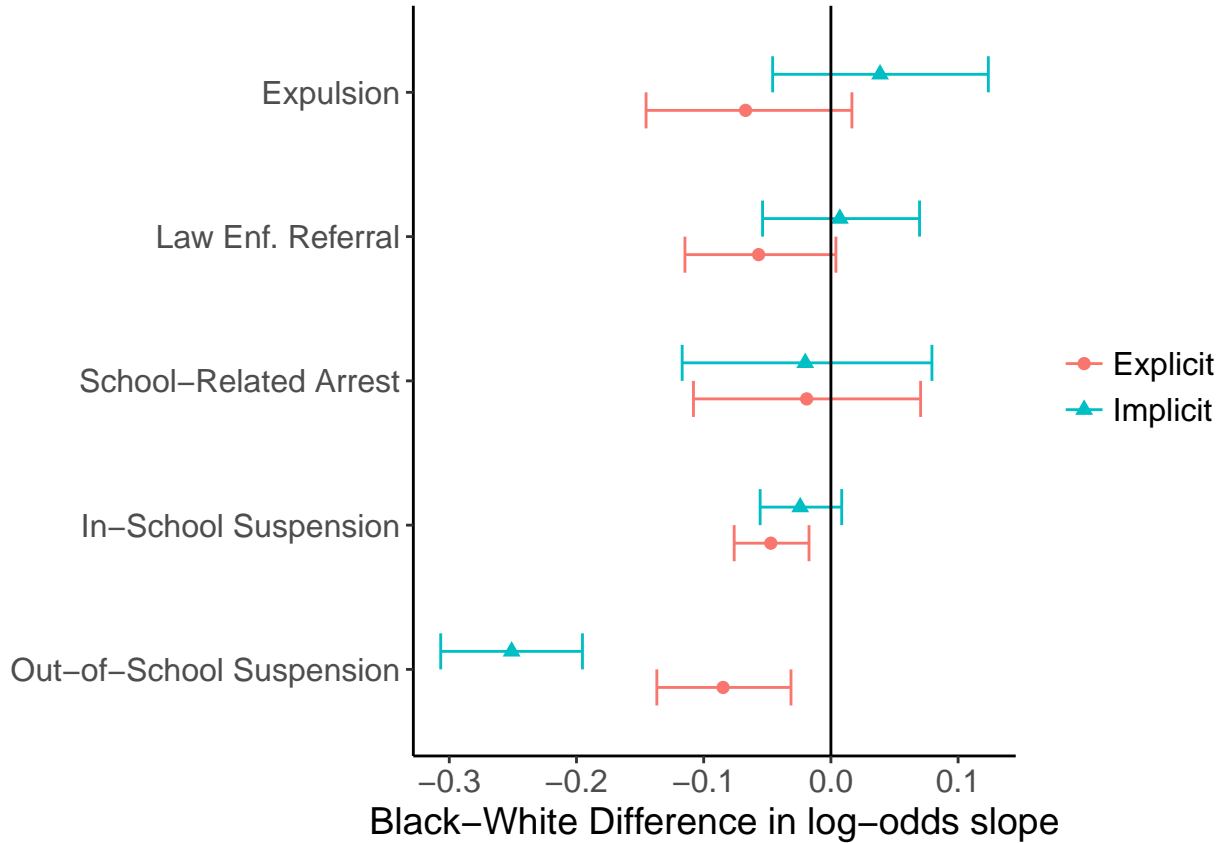


Figure 1. Association between each metric and county-level estimates of explicit and implicit bias. Negative values indicate that the rate of increase (or decrease) for blacks is faster (or slower) than for whites. Point is the mean of the posterior and error bars represent 95% bayesian uncertainty intervals.

estimated at -0.02 $[-0.06, 0.01]$ $p_{neg} = 0.93$, and for explicit bias, the effect is slightly larger, and a positive effect is essentially not credible, given the data and model -0.05 $[-0.08, -0.02]$ $p_{neg} > .99$.

Other outcomes are estimated with less precision, or with patterns that are inconsistent between implicit and explicit bias. For instance, examining the associations for expulsions, the effect of explicit biases are generally in the expected direction ($est = -0.07$, $[-0.15, 0.02]$, $p_{neg} = .94$), but the effect for implicit bias is estimated to be close to zero, with a enough uncertainty ($est = 0.04$, $[-0.05, 0.12]$, $p_{neg} = .19$) to make it difficult to claim an effect of one direction or the other. The model indicates similar uncertainty with respect to

181 school-related arrests and both estimates of bias ($est_{explicit} = -0.02$, $[-0.11, 0.07]$, $p_{neg} = .67$;
 182 $est_{implicit} = -0.02$, $[-0.12, 0.08]$, $p_{neg} = .66$).

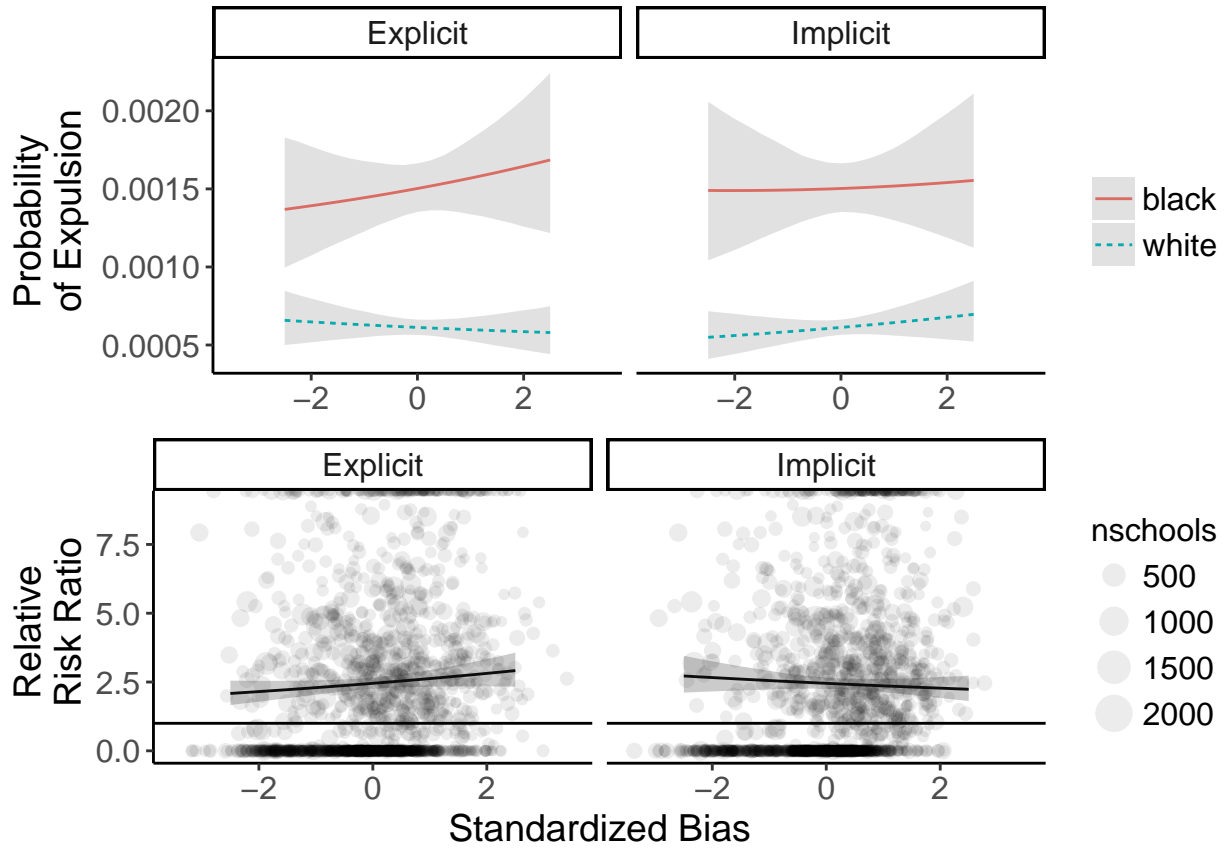


Figure 2. Association between bias and expulsions. Top: Association between bias and the estimated probability of expulsion. Line is the mean of the posterior. Bands indicate 95% uncertainty intervals; Bottom: Association between bias and the relative risk ratio for black students to white students. Points represent counties, whose size are scaled to the number of schools in that county.

183 To better illustrate the nature of these relationships, figure 2 shows the estimated
 184 probabilities of expulsion for black and white students as a function of bias, along with the
 185 relative risk ratio for black students. The relative risk ratio is the ratio of the probability
 186 that a black student will be expelled to the probability that a white student will get expelled.
 187 Values over 1 reflect higher levels of punishment for black students. As previously indicated
 188 in table 1, the top part of this figure makes plain the higher probability of expulsion for

black children. In a county at the mean of the distribution of bias, approximately 0.15% of black students are expected to be expelled [0.14, 0.17]. The corresponding rate for white students is much lower, with about 0.06% expected to be expelled [0.06, 0.07]. Moving to a county one standard deviation above the mean of explicit bias has the effect of increasing the estimated percentage of black students expected to be expelled to 0.16% [0.13, 0.18], while the percentage of white students expected to be expelled would decline to 0.06% [0.05, 0.07]. The same movement for implicit bias would slightly increase the expected expulsions for black students (0.15% [0.13, 0.18]), and increase the expected expulsions for white students a very small amount more to (0.06% [0.06, 0.07]). In real terms, in a county at the mean of the distributions of bias, for every white student expelled, we should expect 2.45 [2.25, 2.45] black students to be expelled. If we move to a county one standard deviation above the mean of explicit bias, the ratio of black to white students expelled increases to 2.63 [2.33, 2.63], while the same movement for implicit bias slightly decreases the ratio to 2.36 [2.13, 2.36].

Figure 3 shows similar patterns, but for out-of-school suspensions. In a county at the mean of the distribution of bias, approximately 11.50% of black students are expected to be suspended [11, 12]. The corresponding rate for white students is just over half that for black students, with about 6.40% expected to be suspended [6.20, 6.50]. Moving to a county one standard deviation above the mean of explicit bias has the effect of slightly decreasing the estimated percentage of black students expected to be suspended to 11.10% [10.30, 11.80], while the percentage of white students expected to be suspended would decrease at a faster rate to 5.60% [5.40, 5.90]. The same movement for implicit bias would dramatically increase the expected suspensions for black students (15.40% [14.40, 16.40]), and increase the expected suspensions for white students a much smaller amount to (6.90% [6.60, 7.30]). In real terms, in a county at the mean of the distributions of bias, for every white student expelled, we should expect 1.80 [1.73, 1.80] black students to be expelled. If we move to a county one standard deviation above the mean of explicit bias, the ratio of black to white students expelled increases to 1.96 [1.84, 1.96], while the same movement for implicit bias

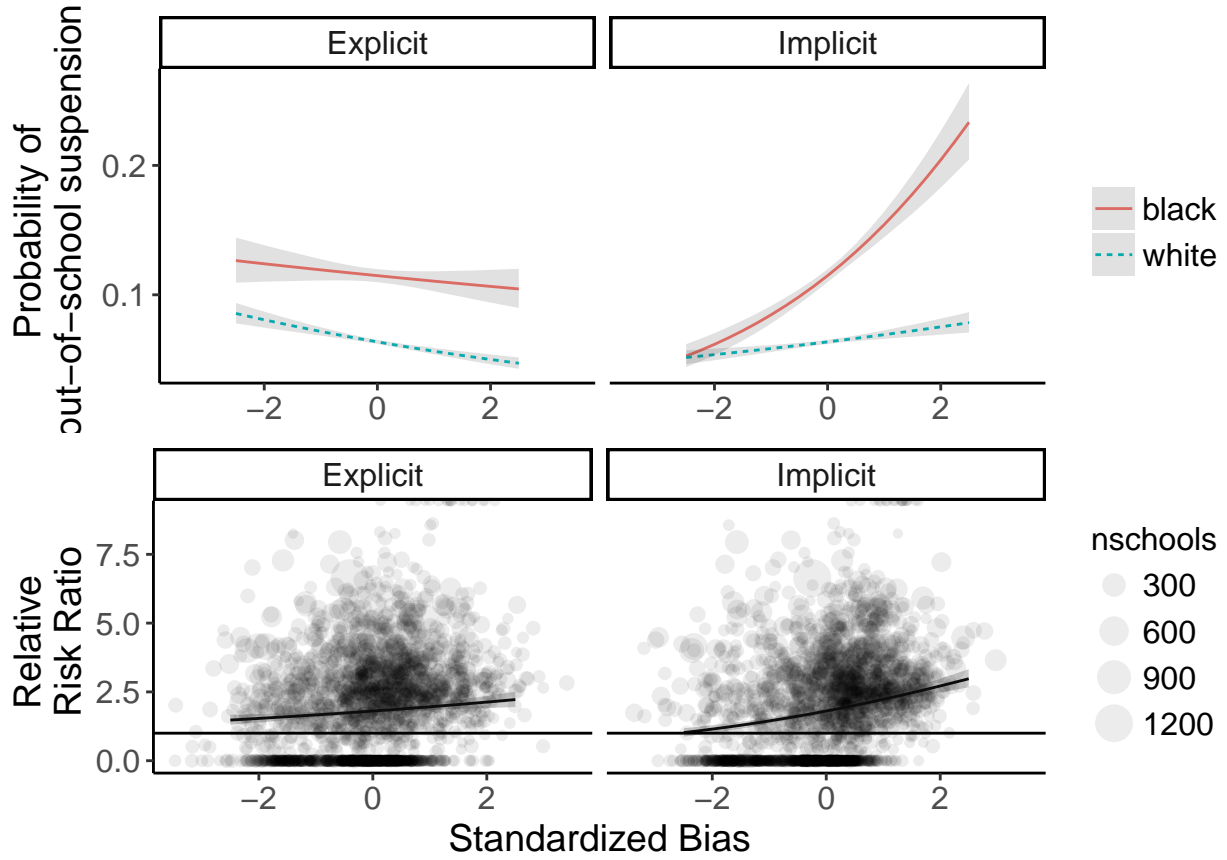


Figure 3. Association between bias and out-of-school suspensions Top: Association between bias and the estimated probability of suspension Line is the mean of the posterior. Bands indicate 95% uncertainty intervals; Bottom: Association between bias and the relative risk ratio for black students to white students. Points represent counties, whose size are scaled to the number of schools in that county.

216 slightly decreases the ratio to 2.23 [2.11, 2.23].⁴

217

Discussion

References

Appendix

Preregistered analysis

In our preanalysis plan, we specified our analyses to focus on 13 actions - corporal punishment, in-school suspension, out-of-school suspension, expulsion with educational services, expulsion without educational services, expulsion under zero-tolerance policies, referral to law enforcement, school-related arrests, mechanical restraint, physical restraint, seclusion, preschool suspension, and preschool expulsion. However, upon further study, we discovered reasons we thought justified excluding a number of these outcomes. In particular, seclusion, physical restraint, and mechanical restraint are not disciplinary actions, but are rather used as means to restrain students who are at risk of harming themselves or others. Additionally, the number of preschool students who are expelled or suspended is vanishingly small (131 total expulsions and 6751 total suspensions out of over 1.4 million enrolled preschool students), making reliably estimating any association across counties exceedingly unlikely. We additionally discovered that counts of one expulsion category (expulsion under zero-tolerance policies) overlapped with counts in other categories, and so excluded this category. To remain consistent with previous studies, we opted to combine the remaining two expulsion categories to yield one overall count of the number of students expelled.

We also preregistered our explicit bias as a simple feeling thermometer towards blacks (i.e. *how warm or cold do you feel towards Blacks? 0=very cold, 10=very warm*). However, past research (Hehman, Flake, & Calanchini, 2017; Leitner et al., 2016) has used the difference in reported warmth towards whites and blacks, and so in the main text, we report models using this metric of explicit bias. Additionally, we preregistered analyses with poststratified estimates (as presented in the main text) along with raw, county-based means. Finally, we had not known about the issues with juvenile justice facilities, or with the school districts with reporting errors. Here, we present the results of the preregistered analyses exactly.

Simple county means

Post stratified estimates

Teacher analyses

- Bates, D., Mächler, M., Bolker, B., & Walker, S. (2015). Fitting linear mixed-effects models using lme4. *Journal of Statistical Software*, 67(1), 1–48. doi:[10.18637/jss.v067.i01](https://doi.org/10.18637/jss.v067.i01)
- Gelman, A., & Little, T. C. (1997). Poststratification into many categories using hierarchical logistic regression. *Survey Methodology*, 23(2), 127–35.
- Gonsalkorale, K., Sherman, J. W., & Klauer, K. C. (2009). Aging and prejudice: Diminished regulation of automatic race bias among older adults. *Journal of Experimental Social Psychology*, 45(2), 410–414.
- Hehman, E., Flake, J. K., & Calanchini, J. (2017). Disproportionate use of lethal force in policing is associated with regional racial biases of residents. *Social Psychological and Personality Science*, 1948550617711229.
- Leitner, J. B., Hehman, E., Ayduk, O., & Mendoza-Denton, R. (2016). Blacks’ death rate due to circulatory diseases is positively related to whites’ explicit racial bias: A nationwide investigation using project implicit. *Psychological Science*, 27(10), 1299–1311.
- Losen, D. J., Hodson, C. L., Keith, I., Michael, A., Morrison, K., & Belway, S. (2015). Are we closing the school discipline gap?
- Miroshnikov, A., & Conlon, E. (2014). *ParallelMCMCcombine: Methods for combining independent subset markov chain monte carlo (mcmc) posterior samples to estimate a posterior density given the full data set*. Retrieved from <https://CRAN.R-project.org/package=parallelMCMCcombine>
- Park, D. K., Gelman, A., & Bafumi, J. (2004). Bayesian multilevel estimation with poststratification: State-level estimates from national polls. *Political Analysis*, 12(4),

375–385.

R Core Team. (2016). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from

<https://www.R-project.org/>

Scott, S. L., Blocker, A. W., Bonassi, F. V., Chipman, H. A., George, E. I., & McCulloch, R. E. (2016). Bayes and big data: The consensus monte carlo algorithm. *International Journal of Management Science and Engineering Management*, 11(2), 78–88.

Stan Development Team. (2016). Rstanarm: Bayesian applied regression modeling via Stan. Retrieved from <http://mc-stan.org/>

Wickham, H. (2009). *Ggplot2: Elegant graphics for data analysis*. Springer-Verlag New York. Retrieved from <http://ggplot2.org>

Wickham, H., & Henry, L. (2017). *Tidyr: Easily tidy data with 'spread()' and 'gather()' functions*. Retrieved from <https://CRAN.R-project.org/package=tidyr>

Wickham, H., Francois, R., Henry, L., & Müller, K. (2017). *Dplyr: A grammar of data manipulation*. Retrieved from <https://CRAN.R-project.org/package=dplyr>

Xu, K., Nosek, B., & Greenwald, A. (2014). Psychology data from the race implicit association test on the project implicit demo website. *Journal of Open Psychology Data*, 2(1).